

Deep Learning with S-shaped Rectified Linear Activation Units

Andrea Lombardo
Luca Tomei

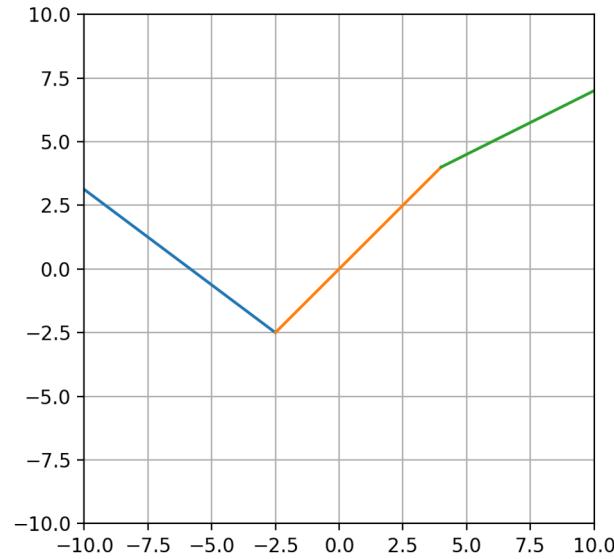
USED ACTIVATION FUNCTIONS

- **RELU:** $\max(0, x_i)$
- **PARAMETRIC RELU:** $\min(0, ax_i) + \max(0, x_i)$
- **LEAKY RELU :** $\min(0, 0.01x_i) + \max(0, x_i)$
- **ELU:** $f(x) = \begin{cases} a(e^x - 1) & \forall x \leq 0 \\ x & \forall x \geq 0 \end{cases}$
- **PARAMETRIC ELU:** $f(x) = \begin{cases} a \left(e^{\frac{x}{b}} - 1 \right) & \forall x \leq 0 \\ \frac{a}{b}x & \forall x \geq 0 \end{cases}$
- **SRELU**

INTRODUCTION OF THE SRELU FUNCTION

S-shaped Rectified Linear Unit

$$f(x_i) = \begin{cases} t_i^l + a_i^l(x_i - t_i^l) & \forall x_i \leq t_i^l \\ x_i & \forall t_i^l < x_i < t_i^r \\ t_i^r + a_i^r(x_i - t_i^r) & \forall x_i \geq t_i^r \end{cases}$$



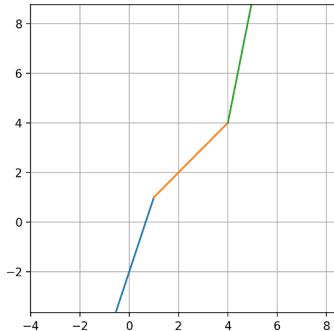
PROPERTIES OF THE SReLU FUNCTION

The SReLU activation function has some remarkable properties:

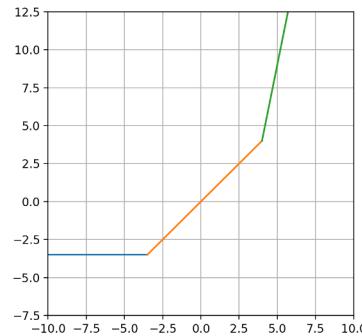
It has four trainable parameters:

$$a_i^r, t_i^r, a_i^l, t_i^l$$

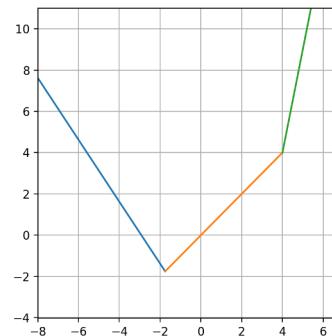
$$a_l > 0$$



$$a_l = 0$$



$$a_l < 0$$

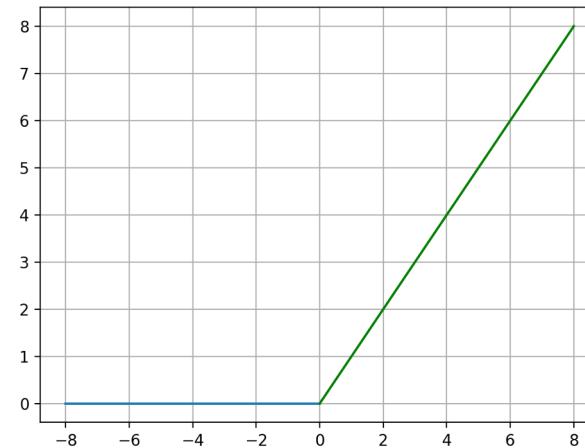


PROPERTIES OF THE SReLU FUNCTION

The SReLU activation function has some remarkable properties:

It generalizes the ReLU function:

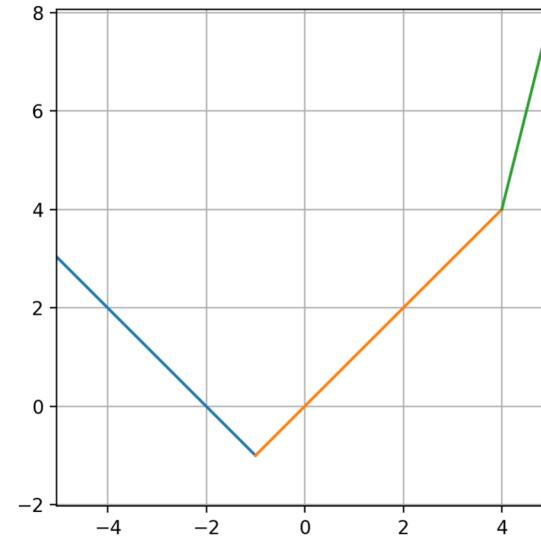
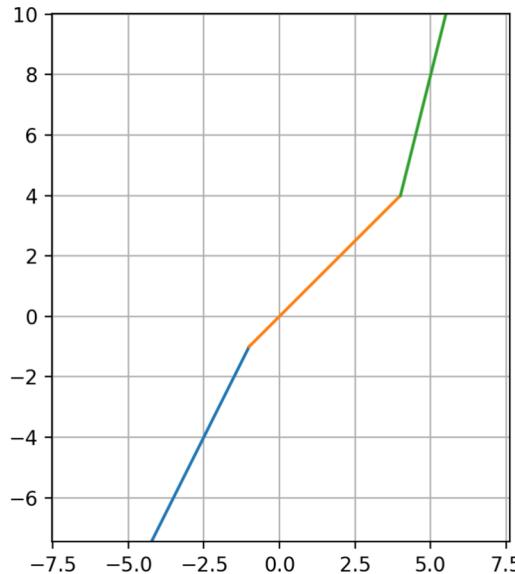
$$\begin{cases} t_l = t_r = 0 \\ a_l = 0 \\ a_r = 1 \end{cases} \implies \text{SReLU} = \text{ReLU}$$



PROPERTIES OF THE SReLU FUNCTION

The SReLU activation function has some remarkable properties:

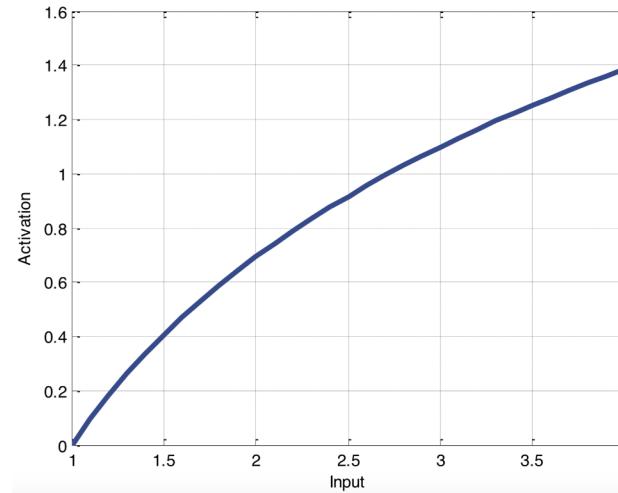
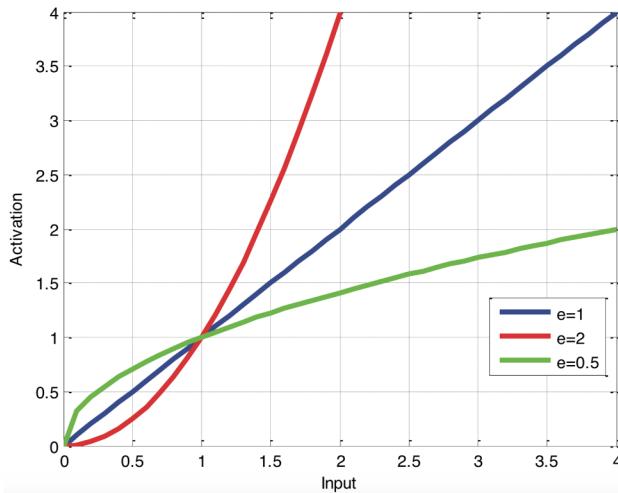
It can be settled to be either convex or non-convex:



PROPERTIES OF THE SReLU FUNCTION

The SReLU activation function has some remarkable properties:

It can be settled to emulate the human perception mechanisms:



EXPERIMENTAL SETTINGS

IMPLEMENTED ARCHITECTURES ON CIFAR10

CNN 1

Model: "sequential"		
Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
activation (Activation)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 30, 30, 32)	9248
activation_1 (Activation)	(None, 30, 30, 32)	0
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
dropout (Dropout)	(None, 15, 15, 32)	0
conv2d_2 (Conv2D)	(None, 15, 15, 64)	18496
activation_2 (Activation)	(None, 15, 15, 64)	0
conv2d_3 (Conv2D)	(None, 13, 13, 64)	36928
activation_3 (Activation)	(None, 13, 13, 64)	0
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
dropout_1 (Dropout)	(None, 6, 6, 64)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 512)	1180160
activation_4 (Activation)	(None, 512)	0
dropout_2 (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 10)	5130
activation_5 (Activation)	(None, 10)	0
<hr/>		
Total params:	1,258,858	
Trainable params:	1,250,858	
Non-trainable params:	0	

CNN 2

Model: "sequential"		
Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
activation (Activation)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
activation_1 (Activation)	(None, 32, 32, 32)	0
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
activation_2 (Activation)	(None, 16, 16, 64)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
activation_3 (Activation)	(None, 16, 16, 64)	0
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
activation_4 (Activation)	(None, 8, 8, 128)	0
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
activation_5 (Activation)	(None, 8, 8, 128)	0
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 128)	262272
activation_6 (Activation)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1290
activation_7 (Activation)	(None, 10)	0
<hr/>		
Total params:	550,570	
Trainable params:	550,570	
Non-trainable params:	0	

IMPLEMENTED ARCHITECTURES ON MNIST

CNN 1

Model: "sequential_1"		
Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 26, 26, 32)	320
activation_1 (Activation)	(None, 26, 26, 32)	0
conv2d_3 (Conv2D)	(None, 24, 24, 64)	18496
activation_2 (Activation)	(None, 24, 24, 64)	0
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_2 (Dropout)	(None, 12, 12, 64)	0
flatten_1 (Flatten)	(None, 9216)	0
dense_2 (Dense)	(None, 128)	1179776
activation_3 (Activation)	(None, 128)	0
dropout_3 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1290
activation_4 (Activation)	(None, 10)	0

Total params: 1,199,882
Trainable params: 1,199,882
Non-trainable params: 0

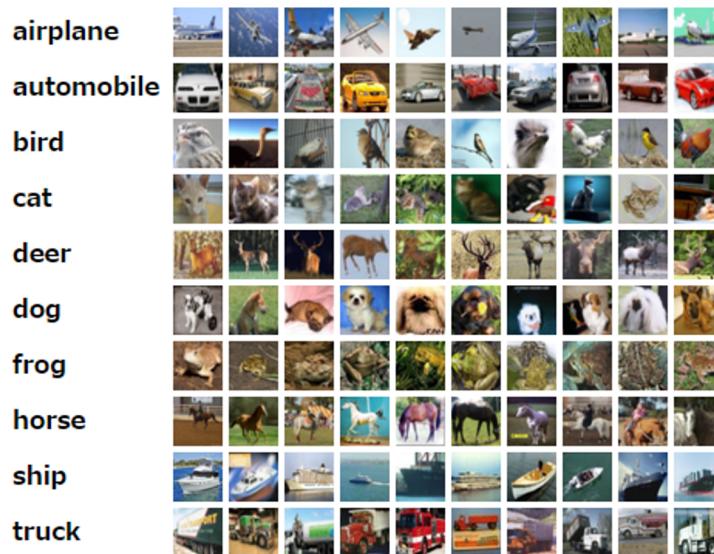
CNN 2

Model: "sequential"		
Layer (type)	output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
activation (Activation)	(None, 26, 26, 32)	0
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
flatten (Flatten)	(None, 5408)	0
dense (Dense)	(None, 100)	540900
activation_1 (Activation)	(None, 100)	0
dense_1 (Dense)	(None, 10)	1010

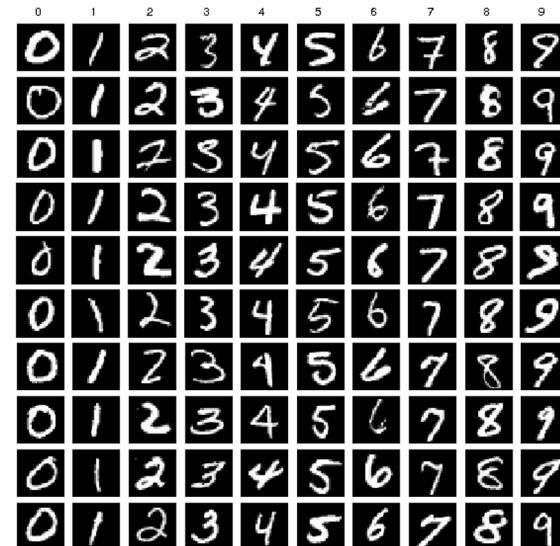
Total params: 542,230
Trainable params: 542,230
Non-trainable params: 0

DATASETS

CIFAR10



MNIST



IMPLEMENTATION DETAILS

NEURAL NETWORKS IMPLEMENTATION

CNN1 on CIFAR10

```
1 model = Sequential()
2 model.add(Conv2D(32, (3, 3), padding='same', input_shape=x_train.shape[1:]))
3 #model.add(Activation('relu'))
4 model.add(LeakyReLU())
5
6 model.add(Conv2D(32, (3, 3)))
7 #model.add(Activation('relu'))
8 model.add(LeakyReLU())
9
10 model.add(MaxPooling2D(pool_size=(2, 2)))
11 model.add(Dropout(0.25))
12 model.add(Conv2D(64, (3, 3), padding='same'))
13 #model.add(Activation('relu'))
14 model.add(LeakyReLU())
15
16 model.add(Conv2D(64, (3, 3)))
17 #model.add(Activation('relu'))
18 model.add(LeakyReLU())
19
20 model.add(MaxPooling2D(pool_size=(2, 2)))
21 model.add(Dropout(0.25))
22
23 model.add(Flatten())
24
25 model.add(Dense(512))
26 #model.add(Activation('relu'))
27 model.add(LeakyReLU())
28
29 model.add(Dense(num_classes))
30 #model.add(Activation('relu'))
31 model.add(LeakyReLU())
32
33 model.add(Dropout(0.5))
34 model.add(Dense(num_classes))
35 model.add(Activation('softmax'))
```

CNN1 on MNIST

```
1 model = Sequential()
2 model.add(Conv2D(32, kernel_size=(3, 3), input_shape=input_shape))
3
4 #model.add(Activation('relu'))
5 model.add(LeakyReLU())
6
7 model.add(Conv2D(64, (3, 3)))
8
9 #model.add(Activation('relu'))
10 model.add(LeakyReLU())
11
12 model.add(MaxPooling2D(pool_size=(2, 2)))
13 model.add(Dropout(0.25))
14 model.add(Flatten())
15 model.add(Dense(128))
16
17 #model.add(Activation('relu'))
18 model.add(LeakyReLU())
19
20 model.add(Dropout(0.5))
21 model.add(Dense(num_classes))
22 model.add(Activation('softmax'))
```

NEURAL NETWORKS IMPLEMENTATION

CNN2 on CIFAR10

```
1 model = Sequential()
2
3 model.add(Conv2D(32, (nb_conv, nb_conv), padding='same', input_shape=x_train.shape[1:]))
4 model.add(Activation('relu'))
5
6 model.add(Conv2D(32, (nb_conv, nb_conv), padding='same'))
7 model.add(Activation('relu'))
8
9 model.add(MaxPooling2D((nb_pool, nb_pool)))
10
11 model.add(Conv2D(64, (nb_conv, nb_conv), padding='same'))
12 model.add(Activation('relu'))
13
14
15 model.add(Conv2D(64, (nb_conv, nb_conv), padding='same'))
16 model.add(Activation('relu'))
17
18 model.add(MaxPooling2D((nb_pool, nb_pool)))
19
20 model.add(Conv2D(128, (nb_conv, nb_conv), padding='same'))
21 model.add(Activation('relu'))
22
23 model.add(Conv2D(128, (nb_conv, nb_conv), padding='same'))
24 model.add(Activation('relu'))
25
26 model.add(MaxPooling2D((nb_pool, nb_pool)))
27
28 model.add(Flatten())
29 model.add(Dense(128))
30 model.add(Activation('relu'))
31
32 model.add(Dense(num_classes))
33 model.add(Activation('softmax'))
```

CNN2 on MNIST

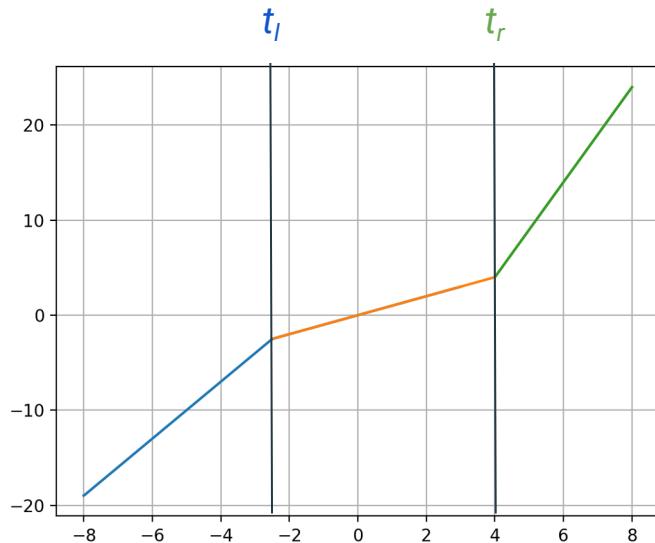
```
1 model = Sequential()
2 model.add(Conv2D(nb_filters, (nb_conv, nb_conv), input_shape=input_shape)
3 model.add(Activation('relu'))
4 #model.add(LeakyReLU())
5
6 model.add(MaxPooling2D((nb_pool, nb_pool)))
7 model.add(Flatten())
8
9 model.add(Dense(100))
10 model.add(Activation('relu'))
11 #model.add(LeakyReLU())
12
13 model.add(Dense(num_classes, activation='softmax'))
```

MySReLU IMPLEMENTATION

```
1  from keras import backend as K
2  from keras.layers import Layer
3
4  class MySReLU(Layer):
5
6      def __init__(self, **kwargs):
7          super(MySReLU, self).__init__(**kwargs)
8
9      def build(self, input_shape):
10         param_shape = tuple(list(input_shape[1:])) # input_shape is: (batch, height, width, channels)
11
12         self.tl = self.add_weight(shape=param_shape, name='tl', initializer='zeros', trainable=True)
13         self.al = self.add_weight(shape=param_shape, name='al', initializer='uniform', trainable=True)
14         self.delta = self.add_weight(shape=param_shape, name='delta', initializer='uniform', trainable=True)
15         self.ar = self.add_weight(shape=param_shape, name='ar', initializer='ones', trainable=True)
16
17         super(MySReLU, self).build(input_shape)
18
19     def call(self, x):
20         tr = self.tl + K.abs(self.delta)
21         tl = self.tl
22         al = self.al
23         ar = self.ar
24
25         eps=0.00001
26         if_x_gtr_tr = K.relu(x-tr)/(x-tr+eps);
27         if_x_lss_tl = K.relu(tl-x)/(tl-x+eps);
28         if_x_btwn_tlr = (1-if_x_gtr_tr)*(1-if_x_lss_tl);
29
30         return if_x_gtr_tr*(ar*(x-tr) + tr) + if_x_lss_tl*(al*(x-tl) + tl) + if_x_btwn_tlr*x
31
32     def compute_output_shape(self, input_shape):
33         return input_shape
```

MySReLU IMPLEMENTATION

```
if_x_lss_tl*(al*(x-tl) + tl) + if_x_btwn_tlr*x + if_x_gtr_tr*(ar*(x-tr) + tr)
```



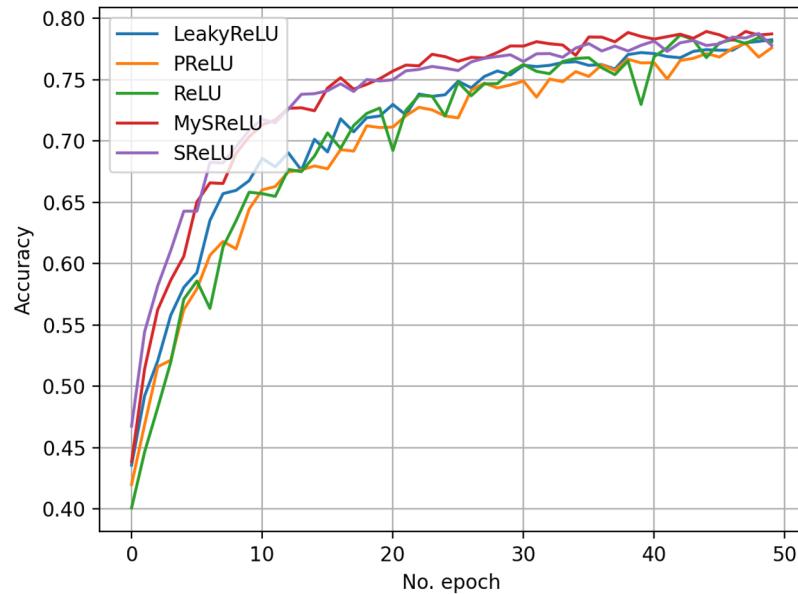
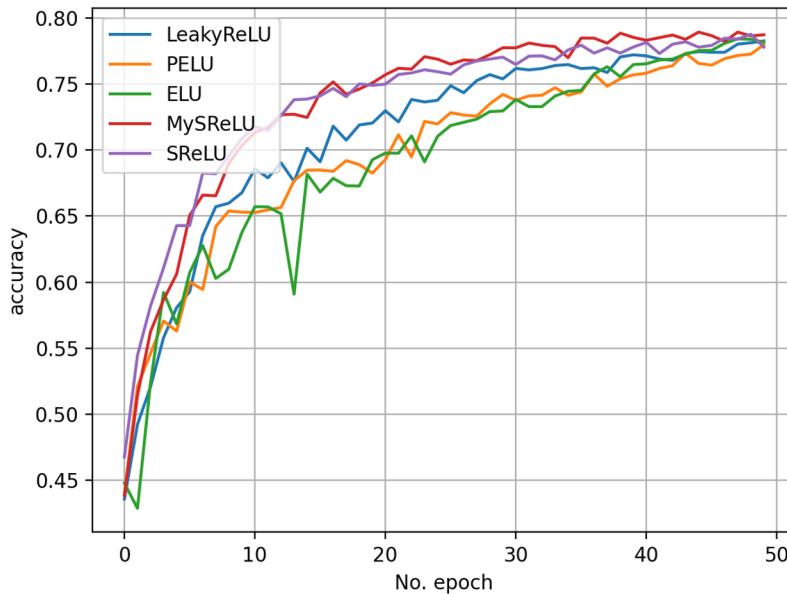
```
if_x_gtr_tr = K.relu(x-tr)/(x-tr+eps)
```

$$\max\left(0, \frac{(x - t_r)}{(x - t_r) + \varepsilon}\right) \simeq \text{Heaviside}(x - t_r)$$

RESULTS

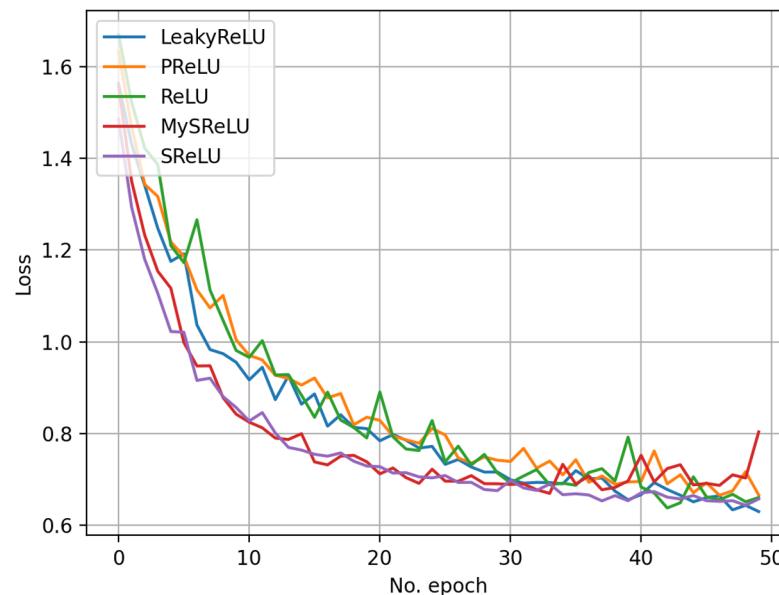
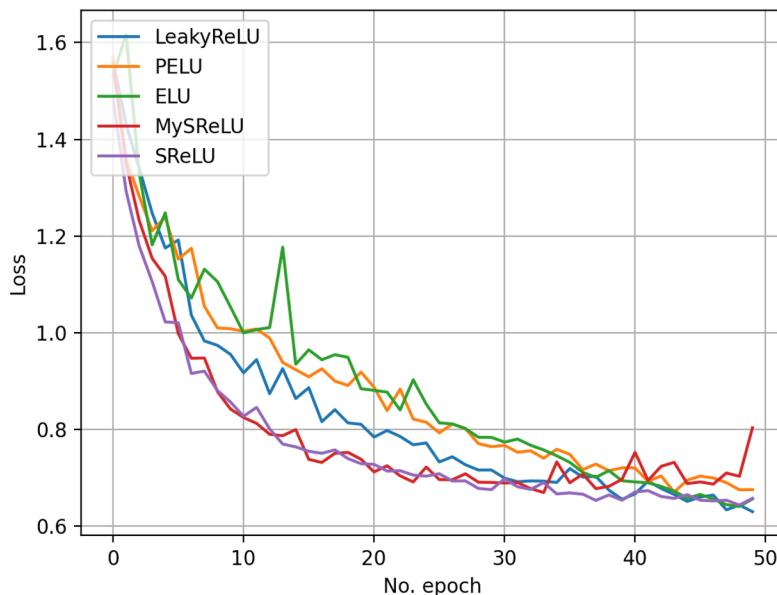
RESULTS: CNN1 ON CIFAR10

ACCURACY: rectified vs exponential activation functions



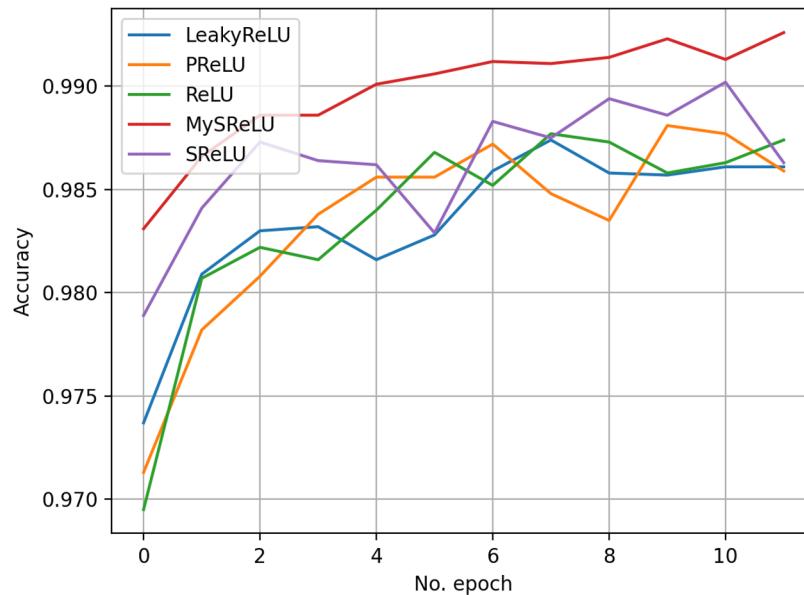
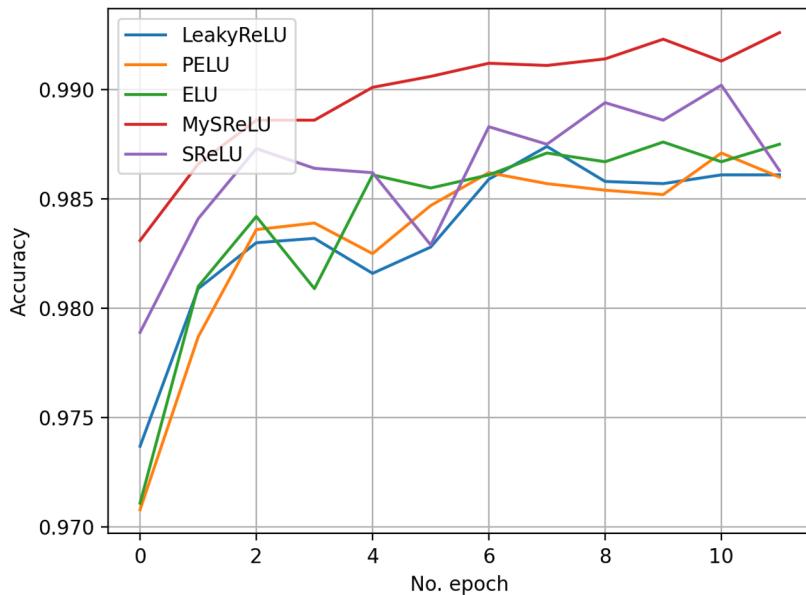
RESULTS: CNN1 ON CIFAR10

LOSS: rectified vs exponential activation functions



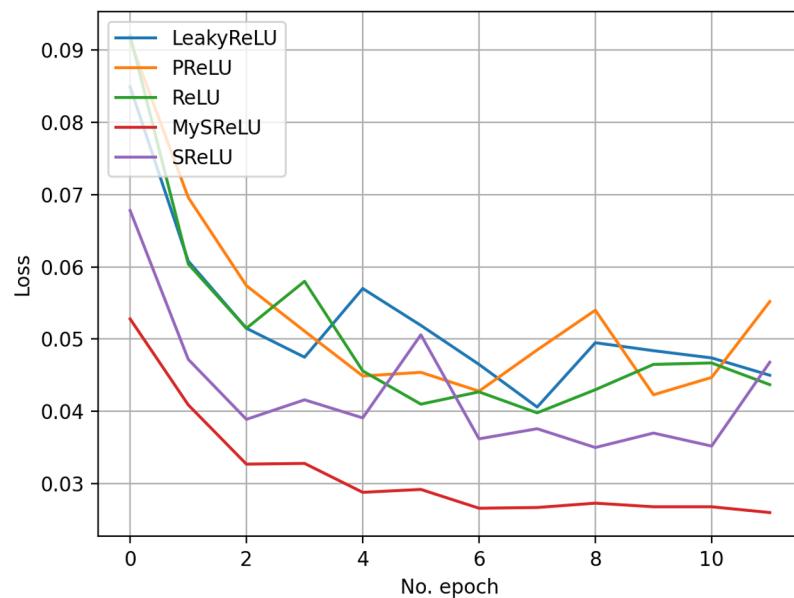
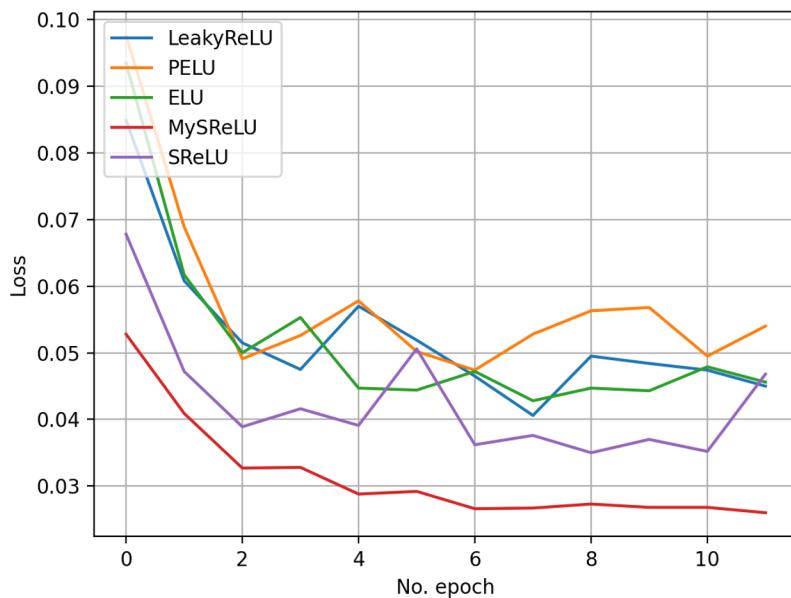
RESULTS: CNN2 ON MNIST

ACCURACY: rectified vs exponential activation functions



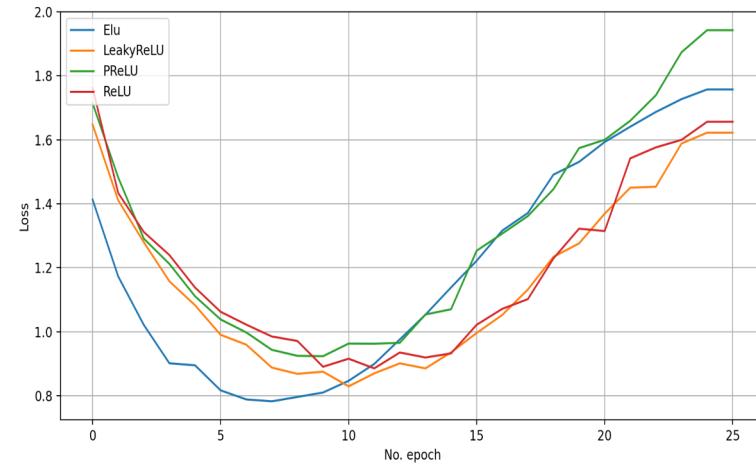
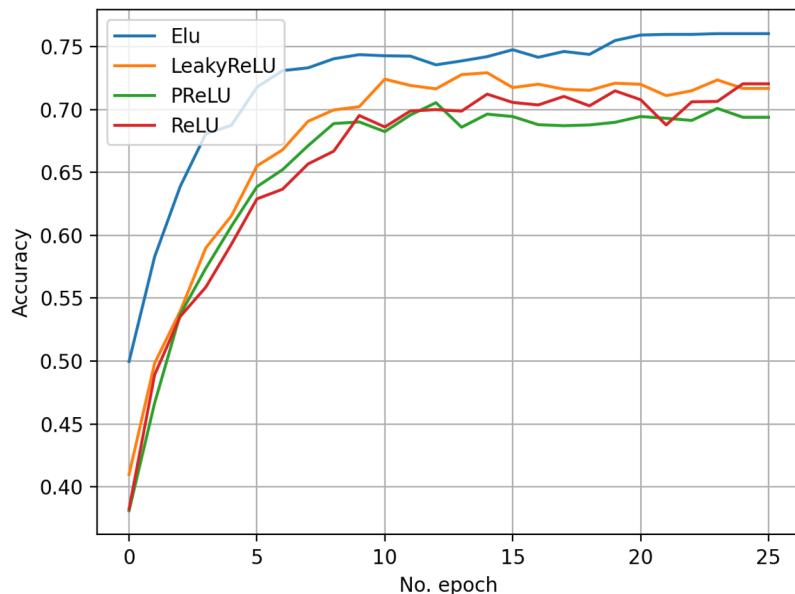
RESULTS: CNN2 ON MNIST

LOSS: rectified vs exponential activation functions



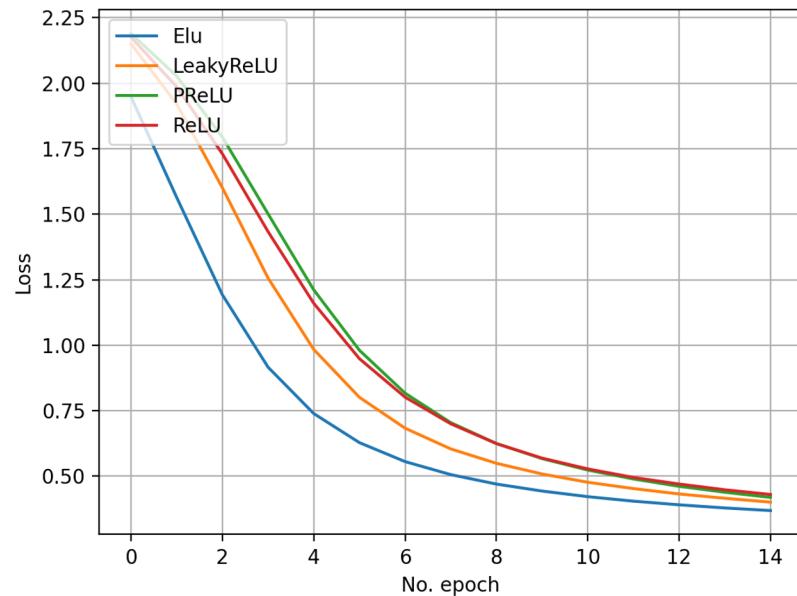
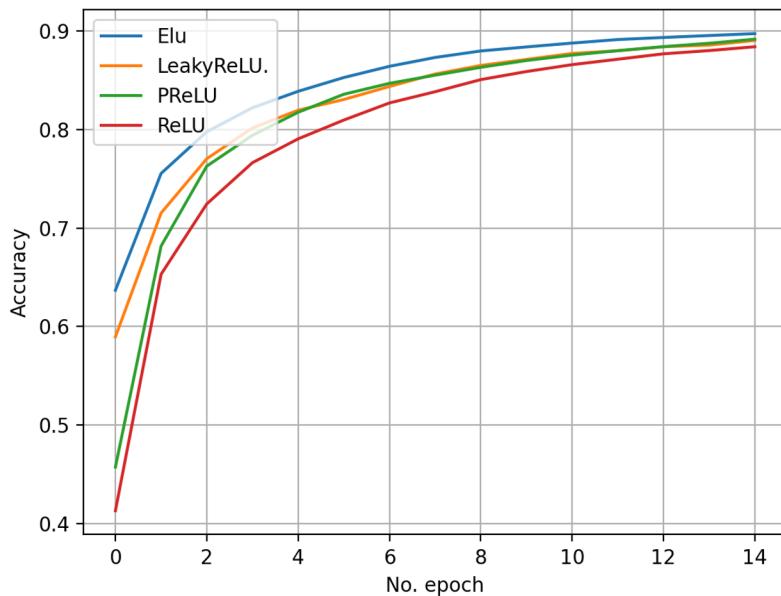
RESULTS: CNN2 ON CIFAR10

ACCURACY & LOSS



RESULTS: CNN1 ON MNIST

ACCURACY & LOSS



Grazie dell'attenzione

